

German Credit Data Exploration_4

Dr. Prashant Mishra

3/27/2018

```
ml_credit_dataset <- read.csv("ml_credit_dataset.csv")
str(ml_credit_dataset)

## 'data.frame':    1000 obs. of  87 variables:
## $ CheckingAccountStatus.0.to.200      : int  0 1 0 0 0 0 0 1 0 1 ...
## $ CheckingAccountStatus.gt.200        : int  0 0 0 0 0 0 0 0 0 0 ...
## $ CheckingAccountStatus.lt.0           : int  1 0 0 1 1 0 0 0 0 0 ...
## $ CheckingAccountStatus.none           : int  0 0 1 0 0 1 1 0 1 0 ...
## $ Duration.0.to.6                     : int  1 0 0 0 0 0 0 0 0 0 ...
## $ Duration.6.to.12                    : int  0 0 1 0 0 0 0 0 1 0 ...
## $ Duration.12.to.18                   : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Duration.18.to.24                   : int  0 0 0 0 1 0 1 0 0 0 ...
## $ Duration.24.to.30                   : int  0 0 0 0 0 0 0 0 0 1 ...
## $ Duration.30.to.36                   : int  0 0 0 0 0 1 0 1 0 0 ...
## $ Duration.36.to.42                   : int  0 0 0 1 0 0 0 0 0 0 ...
## $ Duration.42.to.48                   : int  0 1 0 0 0 0 0 0 0 0 ...
## $ Duration.48.to.54                   : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Duration.54.to.60                   : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Duration.66.to.72                   : int  0 0 0 0 0 0 0 0 0 0 ...
## $ CreditHistory.Critical               : int  1 0 1 0 0 0 0 0 0 1 ...
## $ CreditHistory.Delay                  : int  0 0 0 0 1 0 0 0 0 0 ...
## $ CreditHistory.NoCredit.AllPaid       : int  0 0 0 0 0 0 0 0 0 0 ...
## $ CreditHistory.PaidDuly                : int  0 1 0 1 0 1 1 1 1 0 ...
## $ CreditHistory.ThisBank.AllPaid       : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Purpose.Business                     : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Purpose.DomesticAppliance            : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Purpose.Education                     : int  0 0 1 0 0 1 0 0 0 0 ...
## $ Purpose.Furniture.Equipment          : int  0 0 0 1 0 0 1 0 0 0 ...
## $ Purpose.NewCar                       : int  0 0 0 0 1 0 0 0 0 1 ...
## $ Purpose.Others                       : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Purpose.Radio.Telelevision            : int  1 1 0 0 0 0 0 0 1 0 ...
## $ Purpose.Repairs                       : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Purpose.Retaining                     : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Purpose.UsedCar                       : int  0 0 0 0 0 0 0 1 0 0 ...
## $ SavingsAccountBonds.100.to.500       : int  0 0 0 0 0 0 0 0 0 0 ...
## $ SavingsAccountBonds.500.to.1000     : int  0 0 0 0 0 0 1 0 0 0 ...
## $ SavingsAccountBonds.gt.1000          : int  0 0 0 0 0 0 0 0 1 0 ...
## $ SavingsAccountBonds.lt.100           : int  0 1 1 1 1 0 0 1 0 1 ...
## $ SavingsAccountBonds.Unknown          : int  1 0 0 0 0 1 0 0 0 0 ...
## $ EmploymentDuration.0.to.1            : int  0 0 0 0 0 0 0 0 0 0 ...
## $ EmploymentDuration.1.to.4            : int  0 1 0 0 1 1 0 1 0 0 ...
## $ EmploymentDuration.4.to.7            : int  0 0 1 1 0 0 0 0 1 0 ...
```

```

## $ EmploymentDuration.gt.7 : int 1 0 0 0 0 0 1 0 0 0 ...
## $ EmploymentDuration.Unemployed : int 0 0 0 0 0 0 0 0 0 1 ...
## $ InstallmentRatePercentage.1 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ InstallmentRatePercentage.2 : int 0 1 1 1 0 1 0 1 1 0 ...
## $ InstallmentRatePercentage.3 : int 0 0 0 0 1 0 1 0 0 0 ...
## $ InstallmentRatePercentage.4 : int 1 0 0 0 0 0 0 0 0 1 ...
## $ Personal.Female.NotSingle : int 0 1 0 0 0 0 0 0 0 0 ...
## $ Personal.Male.Divorced.Seperated : int 0 0 0 0 0 0 0 0 1 0 ...
## $ Personal.Male.Married.Widowed : int 0 0 0 0 0 0 0 0 0 1 ...
## $ Personal.Male.Single : int 1 0 1 1 1 1 1 1 0 0 ...
## $ OtherDebtorsGuarantors.CoApplicant : int 0 0 0 0 0 0 0 0 0 0 ...
## $ OtherDebtorsGuarantors.Guarantor : int 0 0 0 1 0 0 0 0 0 0 ...
## $ OtherDebtorsGuarantors.None : int 1 1 1 0 1 1 1 1 1 1 ...
## $ ResidenceDuration.1 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ ResidenceDuration.2 : int 0 1 0 0 0 0 0 1 0 1 ...
## $ ResidenceDuration.3 : int 0 0 1 0 0 0 0 0 0 0 ...
## $ ResidenceDuration.4 : int 1 0 0 1 1 1 1 0 1 0 ...
## $ Property.CarOther : int 0 0 0 0 0 0 0 1 0 1 ...
## $ Property.Insurance : int 0 0 0 1 0 0 1 0 0 0 ...
## $ Property.RealEstate : int 1 1 1 0 0 0 0 0 1 0 ...
## $ Property.Unknown : int 0 0 0 0 1 1 0 0 0 0 ...
## $ Age.18.to.24 : int 0 1 0 0 0 0 0 0 0 0 ...
## $ Age.24.to.30 : int 0 0 0 0 0 0 0 0 0 1 ...
## $ Age.30.to.36 : int 0 0 0 0 0 1 0 1 0 0 ...
## $ Age.36.to.42 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Age.42.to.48 : int 0 0 0 1 0 0 0 0 0 0 ...
## $ Age.48.to.54 : int 0 0 1 0 1 0 1 0 0 0 ...
## $ Age.54.to.60 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Age.60.to.66 : int 0 0 0 0 0 0 0 0 1 0 ...
## $ Age.66.to.72 : int 1 0 0 0 0 0 0 0 0 0 ...
## $ Age.72.to.78 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ OtherInstallmentPlans.Bank : int 0 0 0 0 0 0 0 0 0 0 ...
## $ OtherInstallmentPlans.None : int 1 1 1 1 1 1 1 1 1 1 ...
## $ OtherInstallmentPlans.Stores : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Housing.ForFree : int 0 0 0 1 1 1 0 0 0 0 ...
## $ Housing.Own : int 1 1 1 0 0 0 1 0 1 1 ...
## $ Housing.Rent : int 0 0 0 0 0 0 0 1 0 0 ...
## $ NumberExistingCredits.1 : int 0 1 1 1 0 1 1 1 1 0 ...
## $ NumberExistingCredits.2 : int 1 0 0 0 1 0 0 0 0 1 ...
## $ NumberExistingCredits.3 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ NumberExistingCredits.4 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Job.Management.SelfEmp.HighlyQualified : int 0 0 0 0 0 0 0 1 0 1 ...
## $ Job.SkilledEmployee : int 1 1 0 1 1 0 1 0 0 0 ...
## $ Job.UnemployedUnskilled : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Job.UnskilledResident : int 0 0 1 0 0 1 0 0 1 0 ...
## $ NumberPeopleMaintenance : int 1 1 2 2 2 2 1 1 1 1 ...
## $ Telephone : int 1 0 0 0 0 1 0 1 0 0 ...
## $ ForeignWorker : int 1 1 1 1 1 1 1 1 1 1 ...
## $ Class : Factor w/ 2 levels
"Bad", "Good": 2 1 2 2 1 2 2 2 1 ...

```

Making a Machine Learning task using mlr

```
library(mlr)

## Loading required package: ParamHelpers

credit.task = makeClassifTask(data = ml_credit_dataset, target = "Class")
credit.task = removeConstantFeatures(credit.task)
credit.task

## Supervised task: ml_credit_dataset
## Type: classif
## Target: Class
## Observations: 1000
## Features:
##   numerics      factors    ordered functionals
##       86           0         0             0
## Missings: FALSE
## Has weights: FALSE
## Has blocking: FALSE
## Has coordinates: FALSE
## Classes: 2
##   Bad Good
##   300  700
## Positive class: Bad
```

Cost-sensitive classification

-In regular classification the aim is to minimize the misclassification rate and thus all types of misclassification errors are deemed equally severe.

-A more general setting is cost-sensitive classification where the costs caused by different kinds of errors are not assumed to be equal and the objective is to minimize the expected costs.

-In case of class-dependent costs the costs depend on the true and predicted class label. The costs $c(k,l)$ for predicting class k if the true label is l are usually organized into a $K \times K$ cost matrix where K is the number of classes.

-Naturally, it is assumed that the cost of predicting the correct class label y is minimal (that is $c(y,y) \leq c(k,y)$ for all $k=1,...,K$).

Class-dependent misclassification costs

-There are some classification methods that can accomodate misclassification costs directly. One example is rpart.

-Alternatively, we can use cost-insensitive methods and manipulate the predictions or the training data in order to take misclassification costs into account. mlr supports **thresholding** and **rebalancing**.

-**Thresholding**: The thresholds used to turn posterior probabilities into class labels, are chosen such that the costs are minimized. This requires a Learner that can predict posterior probabilities. During training the costs are not taken into account.

-**Rebalancing**: The idea is to change the proportion of the classes in the training data set in order to account for costs during training, either by `weighting` or by `sampling`. Rebalancing does not require that the Learner can predict probabilities.

— For weighting we need a Learner that supports class weights or observation weights.

— If the Learner cannot deal with weights the proportion of classes can be changed by over- and undersampling.

Cost Matrix for German Credit Data

```
costs = matrix(c(0, 1, 5, 0), 2)
colnames(costs) = rownames(costs) = getTaskClassLevels(credit.task)
costs
```

	Bad	Good
Bad	0	5
Good	1	0

So, the maximum cost is 5 and minimum 0. We penalize if the true class was “Bad” but the model predicts “Good”.

1. Thresholding

We start by fitting a logistic regression model to the German credit data set and predict posterior probabilities.

```
logisticLrn = makeLearner("classif.multinom", predict.type = "prob")

logisticModel = mlr::train(logisticLrn, credit.task)
```

```
## # weights:  88 (87 variable)
## initial   value 693.147181
## iter   10 value 472.774156
## iter   20 value 445.997827
## iter   30 value 444.374321
## iter   40 value 444.223040
## iter   50 value 444.158378
## iter   60 value 444.117755
## iter   70 value 444.107639
## iter   80 value 444.106620
```

```
## final value 444.106579
## converged

logisticpred = predict(logisticModel, task = credit.task)

logisticpred

## Prediction: 1000 observations
## predict.type: prob
## threshold: Bad=0.50,Good=0.50
## time: 0.01
##   id truth   prob.Bad prob.Good response
## 1  1  Good 0.02001323 0.9799868    Good
## 2  2   Bad 0.74111232 0.2588877    Bad
## 3  3  Good 0.03363280 0.9663672    Good
## 4  4  Good 0.10402736 0.8959726    Good
## 5  5   Bad 0.67594919 0.3240508    Bad
## 6  6  Good 0.18333223 0.8166678    Good
## ... (#rows: 1000, #cols: 5)
```

We also fit the data with C50 algorithm.

```
c50Lrn = makeLearner("classif.C50", predict.type = "prob")
c50Model = mlr::train(c50Lrn, credit.task)
c50pred = predict(c50Model, task = credit.task)
c50pred

## Prediction: 1000 observations
## predict.type: prob
## threshold: Bad=0.50,Good=0.50
## time: 0.16
##   id truth   prob.Bad prob.Good response
## 1  1  Good 0.06571429 0.9342857    Good
## 2  2   Bad 0.88750000 0.1125000    Bad
## 3  3  Good 0.08534799 0.9146520    Good
## 4  4  Good 0.04193549 0.9580645    Good
## 5  5   Bad 0.17916667 0.8208333    Good
## 6  6  Good 0.01666667 0.9833333    Good
## ... (#rows: 1000, #cols: 5)
```

i. Theoretical thresholding

The default thresholds for both classes are 0.5. But according to the cost matrix we should predict class Good only if we are very sure that Good is indeed the correct label. Therefore we should increase the threshold for class Good and decrease the threshold for class Bad.

The theoretical threshold for the positive class in two class case can be calculated from the cost matrix as : $t^* = \frac{c(+1,-1)-c(-1,-1)}{c(+1,-1)-c(+1,+1)+c(-1,+1)-c(-1,-1)}$ This formula comes from the fact that cost of predicting class 1 (given the actual is class 1) must be less than cost of predicting -1.

$P(j = -1|x)c(+1, -1) + P(j = +1|x)c(+1, +1) \leq P(j = -1|x)c(-1, -1) + P(j = +1|x)c(-1, +1)$ if we take $p = P(j = +1|x)$ then a threshold value can be derived from, $(1 - t^*)c(+1, -1) + t^*c(-1, -1) = (1 - t^*)c(-1, -1) + t^*c(-1, +1)$

Theoretical threshold

Calculate the theoretical threshold for the positive class: Since $c(+1,+1)=c(-1,-1)=0$

```
th = costs[2,1]/(costs[2,1] + costs[1,2])
th
```

```
## [1] 0.1666667
```

-you can change thresholds in mlr either before training by using the “predict.threshold” option of makeLearner or after prediction by calling setThreshold on the Prediction object.

-Predict class labels according to the theoretical threshold

```
logisticpred.th = setThreshold(logisticpred, th)
logisticpred.th
```

```
## Prediction: 1000 observations
## predict.type: prob
## threshold: Bad=0.17,Good=0.83
## time: 0.01
##   id truth   prob.Bad prob.Good response
## 1  1  Good 0.02001323 0.9799868    Good
## 2  2  Bad 0.74111232 0.2588877    Bad
## 3  3  Good 0.03363280 0.9663672    Good
## 4  4  Good 0.10402736 0.8959726    Good
## 5  5  Bad 0.67594919 0.3240508    Bad
## 6  6  Good 0.18333223 0.8166678    Bad
## ... (#rows: 1000, #cols: 5)
```

```
c50pred.th = setThreshold(c50pred, th)
c50pred.th
```

```
## Prediction: 1000 observations
## predict.type: prob
## threshold: Bad=0.17,Good=0.83
## time: 0.16
##   id truth   prob.Bad prob.Good response
## 1  1  Good 0.06571429 0.9342857    Good
## 2  2  Bad 0.88750000 0.1125000    Bad
## 3  3  Good 0.08534799 0.9146520    Good
## 4  4  Good 0.04193549 0.9580645    Good
## 5  5  Bad 0.17916667 0.8208333    Bad
## 6  6  Good 0.01666667 0.9833333    Good
## ... (#rows: 1000, #cols: 5)
```

Creating a cost measure

In order to calculate the average costs over the entire data set we first need to create a new performance Measure. This can be done through function `makeCostMeasure`. It is expected that the rows of the cost matrix indicate true and the columns predicted class labels.

```
credit.costs = makeCostMeasure(id = "credit.costs", name = "Credit costs",
costs = costs,
  best = 0, worst = 5)
credit.costs

## Name: Credit costs
## Performance measure: credit.costs
## Properties:
classif,classif.multi,req.pred,req.truth,predtype.response,predtype.prob
## Minimize: TRUE
## Best: 0; Worst: 5
## Aggregated by: test.mean
## Arguments: costs=<matrix>, combine=<function>
## Note:
```

Performance measure : Credit cost and Error

Then the average costs can be computed by function `performance`. Below we compare the average costs and the error rate (mmce) of the learning algorithm with both default thresholds 0.5 and theoretical thresholds.

Performance with default thresholds 0.5

```
performance(logisticpred, measures = list(credit.costs, mmce))

## credit.costs      mmce
##      0.735      0.207

calculateConfusionMatrix(logisticpred, relative = TRUE)

## Relative confusion matrix (normalized by row/column):
##      predicted
## true      Bad      Good      -err.-
## Bad      0.56/0.69 0.44/0.17 0.44
## Good      0.11/0.31 0.89/0.83 0.11
## -err.-      0.31      0.17 0.21
##
##
## Absolute confusion matrix:
##      predicted
## true      Bad Good -err.-
## Bad      168 132 132
## Good       75 625 75
## -err.-     75 132 207
```

```

performance(c50pred, measures = list(credit.costs, mmce))

## credit.costs      mmce
##      0.286      0.078

calculateConfusionMatrix(c50pred, relative = TRUE)

## Relative confusion matrix (normalized by row/column):
##      predicted
## true      Bad      Good      -err.-
## Bad      0.83/0.91 0.17/0.07 0.17
## Good      0.04/0.09 0.96/0.93 0.04
## -err.-      0.09      0.07 0.08
##
##
## Absolute confusion matrix:
##      predicted
## true      Bad Good -err.-
## Bad      248  52   52
## Good      26 674   26
## -err.-    26  52   78

```

Performance with theoretical thresholds

```

performance(logisticpred.th, measures = list(credit.costs, mmce))

## credit.costs      mmce
##      0.459      0.339

calculateConfusionMatrix(logisticpred.th, relative = TRUE)

## Relative confusion matrix (normalized by row/column):
##      predicted
## true      Bad      Good      -err.-
## Bad      0.90/0.47 0.10/0.07 0.10
## Good      0.44/0.53 0.56/0.93 0.44
## -err.-      0.53      0.07 0.34
##
##
## Absolute confusion matrix:
##      predicted
## true      Bad Good -err.-
## Bad      270  30   30
## Good      309 391  309
## -err.-    309  30  339

performance(c50pred.th, measures = list(credit.costs, mmce))

## credit.costs      mmce
##      0.263      0.127

calculateConfusionMatrix(c50pred.th, relative = TRUE)

```



```
## Relative confusion matrix (normalized by row/column):
##           predicted
## true      Bad      Good      -err.-
## Bad      0.89/0.74 0.11/0.05 0.11
## Good      0.13/0.26 0.87/0.95 0.13
## -err.-      0.26      0.05 0.13
##
##
## Absolute confusion matrix:
##           predicted
## true      Bad Good -err.-
## Bad      266  34   34
## Good      93 607   93
## -err.-   93  34  127
```

Getting Performance measure with Cross-Validation

These performance values may be overly optimistic as we used the same data set for training and prediction, and resampling strategies should be preferred.

Cross-validated performance with theoretical thresholds

we create a ResampleInstance (rin) that is used throughout the next several code chunks to get comparable performance values.

```
rin = makeResampleInstance("CV", iters = 5, task = credit.task, stratify=TRUE)
```

```
logisticLrn = makeLearner("classif.multinom", predict.type = "prob",
predict.threshold = th, trace = FALSE)
```

```
logisticR = resample(logisticLrn, credit.task, resampling = rin, measures =
list(credit.costs, mmce), show.info = FALSE)
```

```
logisticR
```

```
## Resample Result
## Task: ml_credit_dataset
## Learner: classif.multinom
## Aggr perf: credit.costs.test.mean=0.5690000, mmce.test.mean=0.3610000
## Runtime: 0.690194
```

```
calculateConfusionMatrix(logisticR$pred)
```

```
##           predicted
## true      Bad Good -err.-
## Bad      248  52   52
## Good      309 391  309
## -err.-   309  52  361
```

```
c50rin = makeResampleInstance("CV", iters = 2, task =
credit.task, stratify=TRUE)
```

```
c50Lrn = makeLearner("classif.C50", predict.type = "prob", predict.threshold = th)
c50R = resample(c50Lrn, credit.task, resampling = c50rin, measures = list(credit.costs, mmce), show.info = FALSE)
c50R
```

```
## Resample Result
## Task: ml_credit_dataset
## Learner: classif.C50
## Aggr perf: credit.costs.test.mean=0.8620000,mmce.test.mean=0.3380000
## Runtime: 0.503042
```

```
calculateConfusionMatrix(c50R$pred)
```

```
##           predicted
## true      Bad Good -err.-
##   Bad    169 131   131
##   Good   207 493   207
## -err.-  207 131   338
```

- If we are also interested in the cross-validated performance for the default threshold values we can call `setThreshold` on the resample prediction `r$pred`.
- Cross-validated performance with default thresholds

```
performance(setThreshold(logisticR$pred, 0.5), measures = list(credit.costs, mmce))
```

```
## credit.costs      mmce
##           0.897      0.261
```

```
calculateConfusionMatrix(setThreshold(logisticR$pred, 0.5))
```

```
##           predicted
## true      Bad Good -err.-
##   Bad    141 159   159
##   Good   102 598   102
## -err.-  102 159   261
```

```
performance(setThreshold(c50R$pred, 0.5), measures = list(credit.costs, mmce))
```

```
## credit.costs      mmce
##           0.997      0.285
```

```
calculateConfusionMatrix(setThreshold(c50R$pred, 0.5))
```

```
##           predicted
## true      Bad Good -err.-
##   Bad    122 178   178
##   Good   107 593   107
## -err.-  107 178   285
```

Theoretical threshold vs Performance


- Theoretical thresholding is only reliable if the predicted posterior probabilities are correct. If there is bias the thresholds have to be shifted accordingly.

-Useful in this regard is function “plotThreshVsPerf” that you can use to plot the average costs as well as any other performance measure versus possible threshold values for the positive class in [0,1]. The underlying data is generated by “generateThreshVsPerfData”.


-The following plots show the cross-validated costs and error rate (mmce). The theoretical threshold th calculated above is indicated by the vertical line. As you can see from the left-hand plot the theoretical threshold seems a bit large.

Vertical line is theoretical threshold value.

```
ld = generateThreshVsPerfData(logisticR, measures = list(fpr, tpr,
  credit.costs, mmce))
plotThreshVsPerf(ld, mark.th = th)
```




```
plotROCCurves(ld)
```




```
performance(logisticR$pred, credit.costs)

## credit.costs
##          0.569

cd = generateThreshVsPerfData(c50R, measures = list(fpr, tpr, credit.costs,
mmce))
plotThreshVsPerf(cd, mark.th = th)
```



```
plotROCCurves(cd)
```



```
performance(c50R$pred, credit.costs)
```

```
## credit.costs
##          0.862
```

Learning Curve for various learners

```
r = generateLearningCurveData(
  learners =
c("classif.multinom", "classif.C50", "classif.randomForest", "classif.binomial",
"classif.naiveBayes", "classif.nnet", "classif.rpart"),
  task = credit.task,
  percs = seq(0.1, 1, by = 0.2),
  measures = list(credit.costs, mmce),
  resampling = rin,
  show.info = FALSE)

## # weights:  88 (87 variable)
## initial  value 55.451774
## iter   10 value 3.807073
## iter   20 value 0.010019
## final   value 0.000052
## converged
```

```
## # weights:  88 (87 variable)
## initial  value 55.451774
## iter   10 value 9.443865
## iter   20 value 0.022137
## final   value 0.000065
## converged
## # weights:  88 (87 variable)
## initial  value 55.451774
## iter   10 value 1.287401
## iter   20 value 0.003192
## final   value 0.000062
## converged
## # weights:  88 (87 variable)
## initial  value 55.451774
## iter   10 value 1.325522
## iter   20 value 0.003499
## final   value 0.000063
## converged
## # weights:  88 (87 variable)
## initial  value 55.451774
## iter   10 value 6.642383
## iter   20 value 0.019643
## final   value 0.000060
## converged
## # weights:  88 (87 variable)
## initial  value 166.355323
## iter   10 value 80.277877
## iter   20 value 74.892167
## iter   30 value 74.706087
## iter   40 value 74.626202
## iter   50 value 74.578432
## iter   60 value 74.572955
## iter   70 value 74.572272
## final   value 74.572262
## converged
## # weights:  88 (87 variable)
## initial  value 166.355323
## iter   10 value 75.517451
## iter   20 value 63.799106
## iter   30 value 62.571838
## iter   40 value 62.249788
## iter   50 value 62.161348
## iter   60 value 62.152748
## final   value 62.152212
## converged
## # weights:  88 (87 variable)
## initial  value 166.355323
## iter   10 value 67.284231
## iter   20 value 54.758069
## iter   30 value 53.991517
```

```
## iter 40 value 53.871575
## iter 50 value 53.850633
## iter 60 value 53.841576
## iter 70 value 53.838375
## iter 80 value 53.835746
## iter 90 value 53.835561
## final value 53.835556
## converged
## # weights: 88 (87 variable)
## initial value 166.355323
## iter 10 value 84.632420
## iter 20 value 77.208149
## iter 30 value 76.828918
## iter 40 value 76.726372
## iter 50 value 76.697820
## iter 60 value 76.691755
## iter 70 value 76.690438
## final value 76.690377
## converged
## # weights: 88 (87 variable)
## initial value 166.355323
## iter 10 value 71.480459
## iter 20 value 63.606216
## iter 30 value 63.209643
## iter 40 value 63.157022
## iter 50 value 63.135414
## iter 60 value 63.129495
## iter 70 value 63.128884
## final value 63.128870
## converged
## # weights: 88 (87 variable)
## initial value 277.258872
## iter 10 value 164.236856
## iter 20 value 156.118323
## iter 30 value 155.074829
## iter 40 value 154.856918
## iter 50 value 154.833365
## iter 60 value 154.830752
## iter 70 value 154.830450
## final value 154.830447
## converged
## # weights: 88 (87 variable)
## initial value 277.258872
## iter 10 value 166.561430
## iter 20 value 157.732328
## iter 30 value 156.392310
## iter 40 value 156.180983
## iter 50 value 156.165601
## iter 60 value 156.163060
## iter 70 value 156.162758
```



```
## final value 156.162753
## converged
## # weights: 88 (87 variable)
## initial value 277.258872
## iter 10 value 151.966617
## iter 20 value 145.394673
## iter 30 value 145.039999
## iter 40 value 144.875023
## iter 50 value 144.851300
## iter 60 value 144.847803
## iter 70 value 144.847360
## final value 144.847355
## converged
## # weights: 88 (87 variable)
## initial value 277.258872
## iter 10 value 149.419065
## iter 20 value 143.081585
## iter 30 value 142.305337
## iter 40 value 142.065624
## iter 50 value 142.043358
## iter 60 value 142.041400
## iter 70 value 142.041087
## final value 142.041082
## converged
## # weights: 88 (87 variable)
## initial value 277.258872
## iter 10 value 175.359757
## iter 20 value 169.793320
## iter 30 value 169.263472
## iter 40 value 169.121848
## iter 50 value 169.111592
## iter 60 value 169.110673
## final value 169.110657
## converged
## # weights: 88 (87 variable)
## initial value 388.162421
## iter 10 value 242.848981
## iter 20 value 234.259877
## iter 30 value 233.274310
## iter 40 value 233.051893
## iter 50 value 232.979714
## iter 60 value 232.972615
## final value 232.972432
## converged
## # weights: 88 (87 variable)
## initial value 388.162421
## iter 10 value 260.438488
## iter 20 value 244.837789
## iter 30 value 243.605784
## iter 40 value 243.354349
```

```
## iter 50 value 243.273317
## iter 60 value 243.263971
## iter 70 value 243.263131
## final value 243.263073
## converged
## # weights: 88 (87 variable)
## initial value 388.162421
## iter 10 value 241.903448
## iter 20 value 233.685801
## iter 30 value 232.687674
## iter 40 value 232.507637
## iter 50 value 232.451467
## iter 60 value 232.445852
## iter 70 value 232.444924
## final value 232.444803
## converged
## # weights: 88 (87 variable)
## initial value 388.162421
## iter 10 value 234.755989
## iter 20 value 228.581537
## iter 30 value 227.610911
## iter 40 value 227.318516
## iter 50 value 227.271090
## iter 60 value 227.260663
## iter 70 value 227.260303
## final value 227.260205
## converged
## # weights: 88 (87 variable)
## initial value 388.162421
## iter 10 value 245.407613
## iter 20 value 239.930179
## iter 30 value 239.519742
## iter 40 value 239.413945
## iter 50 value 239.386533
## iter 60 value 239.383575
## iter 70 value 239.383368
## final value 239.383356
## converged
## # weights: 88 (87 variable)
## initial value 499.065970
## iter 10 value 339.487205
## iter 20 value 315.072806
## iter 30 value 313.330165
## iter 40 value 313.115924
## iter 50 value 313.114169
## final value 313.114106
## converged
## # weights: 88 (87 variable)
## initial value 499.065970
## iter 10 value 334.180547
```

```
## iter 20 value 313.038276
## iter 30 value 311.057154
## iter 40 value 310.654260
## iter 50 value 310.463211
## iter 60 value 310.428355
## iter 70 value 310.424960
## final value 310.424861
## converged
## # weights: 88 (87 variable)
## initial value 499.065970
## iter 10 value 321.595131
## iter 20 value 293.531266
## iter 30 value 291.081185
## iter 40 value 290.884871
## iter 50 value 290.833315
## iter 60 value 290.814264
## iter 70 value 290.810373
## final value 290.810137
## converged
## # weights: 88 (87 variable)
## initial value 499.065970
## iter 10 value 337.305632
## iter 20 value 312.651695
## iter 30 value 309.815725
## iter 40 value 309.481207
## iter 50 value 309.414246
## iter 60 value 309.391794
## iter 70 value 309.388565
## iter 80 value 309.388431
## iter 80 value 309.388428
## iter 80 value 309.388428
## final value 309.388428
## converged
## # weights: 88 (87 variable)
## initial value 499.065970
## iter 10 value 341.962869
## iter 20 value 319.157459
## iter 30 value 316.865340
## iter 40 value 316.651791
## iter 50 value 316.640647
## final value 316.640585
## converged

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
## Warning: glm.fit: algorithm did not converge
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
## Warning: glm.fit: algorithm did not converge
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
## Warning: glm.fit: algorithm did not converge
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
## Warning: glm.fit: algorithm did not converge
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
## Warning: glm.fit: algorithm did not converge
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
## Warning: glm.fit: algorithm did not converge
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## # weights: 265
## initial value 46.222933
## iter 10 value 14.264704
## iter 20 value 5.644244
## iter 30 value 0.204875
## iter 40 value 0.017270
## iter 50 value 0.002323
## iter 60 value 0.000337
## final value 0.000087
## converged
## # weights: 265
## initial value 72.425531
## iter 10 value 44.685219
## iter 20 value 21.822272
## iter 30 value 18.889102
## iter 40 value 17.368947
## iter 50 value 17.301437
## iter 60 value 17.294162
## iter 70 value 17.290528
## iter 80 value 17.289201
## iter 90 value 17.288451
## iter 100 value 17.287582
## final value 17.287582
## stopped after 100 iterations
## # weights: 265
## initial value 49.835598
## iter 10 value 16.710599
```

```
## iter 20 value 9.044953
## iter 30 value 8.702449
## iter 40 value 8.699886
## iter 50 value 8.699724
## final value 8.699706
## converged
## # weights: 265
## initial value 61.520181
## iter 10 value 22.325689
## iter 20 value 12.017392
## iter 30 value 11.911723
## final value 11.910927
## converged
## # weights: 265
## initial value 55.091284
## iter 10 value 32.634114
## iter 20 value 5.705752
## iter 30 value 0.035254
## iter 40 value 0.000875
## final value 0.000097
## converged
## # weights: 265
## initial value 201.947936
## iter 10 value 100.790362
## iter 20 value 47.580523
## iter 30 value 40.164610
## iter 40 value 39.034419
## iter 50 value 38.647607
## iter 60 value 38.547760
## iter 70 value 38.224459
## iter 80 value 38.190700
## iter 90 value 38.131963
## iter 100 value 38.077380
## final value 38.077380
## stopped after 100 iterations
## # weights: 265
## initial value 142.984555
## iter 10 value 73.062650
## iter 20 value 47.894417
## iter 30 value 37.579720
## iter 40 value 36.017468
## iter 50 value 34.835247
## iter 60 value 34.759534
## iter 70 value 33.049426
## iter 80 value 31.959051
## iter 90 value 31.899649
## iter 100 value 30.355795
## final value 30.355795
## stopped after 100 iterations
## # weights: 265
```

```
## initial value 152.962914
## iter 10 value 54.285021
## iter 20 value 30.808785
## iter 30 value 21.892041
## iter 40 value 17.795722
## iter 50 value 13.469473
## iter 60 value 11.918724
## iter 70 value 10.364017
## iter 80 value 10.094707
## iter 90 value 10.043747
## iter 100 value 10.023037
## final value 10.023037
## stopped after 100 iterations
## # weights: 265
## initial value 167.995785
## iter 10 value 85.140979
## iter 20 value 45.063909
## iter 30 value 39.537326
## iter 40 value 38.132644
## iter 50 value 37.026263
## iter 60 value 36.473699
## iter 70 value 35.500198
## iter 80 value 35.414661
## iter 90 value 35.316688
## iter 100 value 34.910731
## final value 34.910731
## stopped after 100 iterations
## # weights: 265
## initial value 157.495367
## iter 10 value 86.305798
## iter 20 value 38.730793
## iter 30 value 29.643201
## iter 40 value 26.430338
## iter 50 value 25.532968
## iter 60 value 25.138708
## iter 70 value 25.029303
## iter 80 value 24.932492
## iter 90 value 24.820971
## iter 100 value 24.728426
## final value 24.728426
## stopped after 100 iterations
## # weights: 265
## initial value 245.180323
## iter 10 value 184.816134
## iter 20 value 144.566172
## iter 30 value 132.714597
## iter 40 value 124.805878
## iter 50 value 123.390509
## iter 60 value 123.364285
## final value 123.363526
```



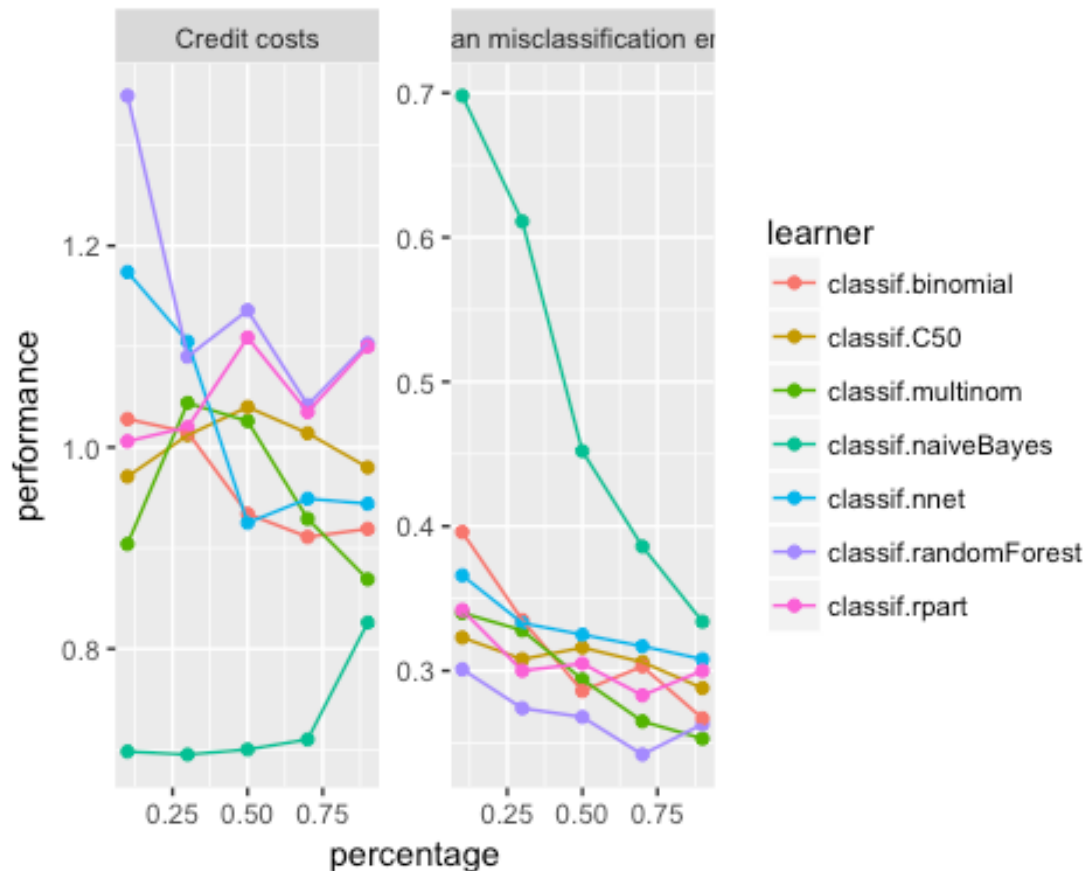
```
## converged
## # weights: 265
## initial value 325.746979
## iter 10 value 167.604994
## iter 20 value 127.113716
## iter 30 value 116.855586
## iter 40 value 114.614752
## iter 50 value 113.248873
## iter 60 value 101.125766
## iter 70 value 71.660513
## iter 80 value 68.422666
## iter 90 value 67.575778
## iter 100 value 65.325190
## final value 65.325190
## stopped after 100 iterations
## # weights: 265
## initial value 261.666995
## iter 10 value 183.638604
## iter 20 value 105.377836
## iter 30 value 76.574531
## iter 40 value 66.753184
## iter 50 value 65.997622
## iter 60 value 65.947301
## iter 70 value 65.942793
## final value 65.942421
## converged
## # weights: 265
## initial value 239.201472
## iter 10 value 152.742680
## iter 20 value 71.554487
## iter 30 value 41.714322
## iter 40 value 29.924505
## iter 50 value 26.608479
## iter 60 value 26.319698
## iter 70 value 26.260016
## iter 80 value 26.215021
## iter 90 value 26.205766
## iter 100 value 26.195195
## final value 26.195195
## stopped after 100 iterations
## # weights: 265
## initial value 341.477202
## iter 10 value 188.550709
## iter 20 value 116.449179
## iter 30 value 95.538472
## iter 40 value 88.159395
## iter 50 value 86.918138
## iter 60 value 86.604530
## iter 70 value 86.598618
## iter 80 value 86.597877
```

```
## iter 90 value 86.597248
## iter 100 value 86.596474
## final value 86.596474
## stopped after 100 iterations
## # weights: 265
## initial value 470.567119
## iter 10 value 240.646144
## iter 20 value 160.131889
## iter 30 value 113.923434
## iter 40 value 96.839468
## iter 50 value 94.604924
## iter 60 value 94.334611
## iter 70 value 94.253879
## iter 80 value 94.211061
## iter 90 value 94.205761
## iter 100 value 94.194246
## final value 94.194246
## stopped after 100 iterations
## # weights: 265
## initial value 353.637329
## iter 10 value 238.725695
## iter 20 value 132.152128
## iter 30 value 96.033282
## iter 40 value 90.584879
## iter 50 value 88.147088
## iter 60 value 87.764506
## iter 70 value 87.550845
## iter 80 value 87.295329
## iter 90 value 87.066890
## iter 100 value 85.488172
## final value 85.488172
## stopped after 100 iterations
## # weights: 265
## initial value 408.940062
## iter 10 value 310.134634
## iter 20 value 239.359800
## iter 30 value 197.571397
## iter 40 value 173.372032
## iter 50 value 160.978936
## iter 60 value 159.475925
## iter 70 value 159.313579
## iter 80 value 149.282052
## iter 90 value 121.670506
## iter 100 value 119.474659
## final value 119.474659
## stopped after 100 iterations
## # weights: 265
## initial value 362.733498
## iter 10 value 237.417465
## iter 20 value 179.617091
```

```
## iter 30 value 152.951412
## iter 40 value 134.837331
## iter 50 value 124.575811
## iter 60 value 119.826501
## iter 70 value 116.702424
## iter 80 value 114.306921
## iter 90 value 112.134618
## iter 100 value 111.875377
## final value 111.875377
## stopped after 100 iterations
## # weights: 265
## initial value 348.756675
## iter 10 value 237.098870
## iter 20 value 182.500195
## iter 30 value 165.231841
## iter 40 value 154.667031
## iter 50 value 153.944042
## iter 60 value 153.724317
## iter 70 value 153.615366
## iter 80 value 153.336663
## iter 90 value 153.296700
## iter 100 value 153.203238
## final value 153.203238
## stopped after 100 iterations
## # weights: 265
## initial value 494.492761
## iter 10 value 344.191345
## iter 20 value 275.217011
## iter 30 value 238.561363
## iter 40 value 213.685137
## iter 50 value 211.459897
## iter 60 value 210.587066
## iter 70 value 210.217026
## iter 80 value 210.198349
## iter 90 value 210.195148
## iter 100 value 210.190061
## final value 210.190061
## stopped after 100 iterations
## # weights: 265
## initial value 506.211030
## iter 10 value 396.185142
## iter 20 value 295.981830
## iter 30 value 252.484851
## iter 40 value 226.589904
## iter 50 value 210.280278
## iter 60 value 193.738521
## iter 70 value 187.886520
## iter 80 value 183.889181
## iter 90 value 183.684237
## iter 100 value 183.674961
```

```
## final value 183.674961
## stopped after 100 iterations
## # weights: 265
## initial value 438.278228
## iter 10 value 322.722620
## iter 20 value 280.754927
## iter 30 value 236.398527
## iter 40 value 205.275631
## iter 50 value 192.216905
## iter 60 value 185.880155
## iter 70 value 176.807511
## iter 80 value 169.670020
## iter 90 value 165.623388
## iter 100 value 165.295149
## final value 165.295149
## stopped after 100 iterations
## # weights: 265
## initial value 518.651516
## iter 10 value 319.563381
## iter 20 value 242.934582
## iter 30 value 217.189300
## iter 40 value 195.554167
## iter 50 value 189.042559
## iter 60 value 187.319527
## iter 70 value 187.304123
## iter 80 value 187.299368
## iter 90 value 187.296706
## iter 100 value 187.295827
## final value 187.295827
## stopped after 100 iterations
## # weights: 265
## initial value 448.928837
## iter 10 value 337.075259
## iter 20 value 302.103383
## iter 30 value 232.626402
## iter 40 value 174.001101
## iter 50 value 151.967871
## iter 60 value 148.625430
## iter 70 value 148.391015
## iter 80 value 148.227692
## iter 90 value 148.138944
## iter 100 value 148.101858
## final value 148.101858
## stopped after 100 iterations
```

```
plotLearningCurve(r)
```



RandomForest model:

```
Randomlrn = makeLearner("classif.randomForest", predict.type = "prob",
fix.factors.prediction = TRUE)
rin = makeResampleInstance("CV", iters = 5, task = credit.task, stratify=TRUE)
Ranr = resample(Randomlrn, credit.task, rin, measures = list(credit.costs,
mmce), show.info = FALSE)
Ranr
```

```
## Resample Result
## Task: ml_credit_dataset
## Learner: classif.randomForest
## Aggr perf: credit.costs.test.mean=1.0190000, mmce.test.mean=0.2510000
## Runtime: 14.9573
```

Prediction based on theoretical threshold

```
Ranpred.th = setThreshold(Ranr$pred, threshold = th)
calculateConfusionMatrix(Ranpred.th)
```

```
##           predicted
## true      Bad Good -err.-
## Bad      280  20   20
```


```
##      Good    430   270    430
##    -err.-  430    20    450

performance(Ranpred.th, measures = list(credit.costs, mmce))


## credit.costs      mmce
##          0.53      0.45
```

Tuning Threshold

```
dr = generateThreshVsPerfData(Ranr, measures = list(fpr, tpr, credit.costs,
mmce))
plotThreshVsPerf(dr, mark.th = th)
```



```
plotROCCurves(dr)
```



```
performance(Ranr$pred, credit.costs)
```

```
## credit.costs
##      1.019
```

Naive Bayes Model:

```
NBlrn = makeLearner("classif.naiveBayes", predict.type = "prob",
  fix.factors.prediction = TRUE)
rin = makeResampleInstance("CV", iters = 5, task = credit.task, stratify=TRUE)
NBr = resample(NBlrn, credit.task, rin, measures = list(credit.costs, mmce),
  show.info = FALSE)
NBr
```

```
## Resample Result
## Task: ml_credit_dataset
## Learner: classif.naiveBayes
## Aggr perf: credit.costs.test.mean=0.8340000, mmce.test.mean=0.3180000
## Runtime: 1.49339
```

```
NBpred.th = setThreshold(NBr$pred, threshold = th)
calculateConfusionMatrix(NBpred.th)
```


```
##           predicted
## true      Bad Good -err.-
## Bad      191 109   109
## Good     229 471   229
## -err.-   229 109   338
```

```
performance(NBpred.th, measures = list(credit.costs, mmce))
```


```
## credit.costs      mmce
##           0.774      0.338
```

```
Nr = generateThreshVsPerfData(NBr, measures = list(fpr, tpr, credit.costs,
mmce))
```

```
plotThreshVsPerf(Nr, mark.th = th)
```



```
plotROCCurves(Nr)
```

```
performance(NBr$pred,credit.costs)
```

```
## credit.costs
##          0.834
```

Binomial Model

```
Blrn = makeLearner("classif.binomial", predict.type = "prob",
  fix.factors.prediction = TRUE)
rin = makeResampleInstance("CV", iters = 5, task = credit.task, stratify=TRUE)
Br = resample(Blrn, credit.task, rin, measures = list(credit.costs, mmce),
  show.info = FALSE)
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
```

```

## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

Br

## Resample Result
## Task: ml_credit_dataset
## Learner: classif.binomial
## Aggr perf: credit.costs.test.mean=0.8810000,mmce.test.mean=0.2610000
## Runtime: 0.766087

Bpred.th = setThreshold(Br$pred, threshold = th)
calculateConfusionMatrix(Bpred.th)

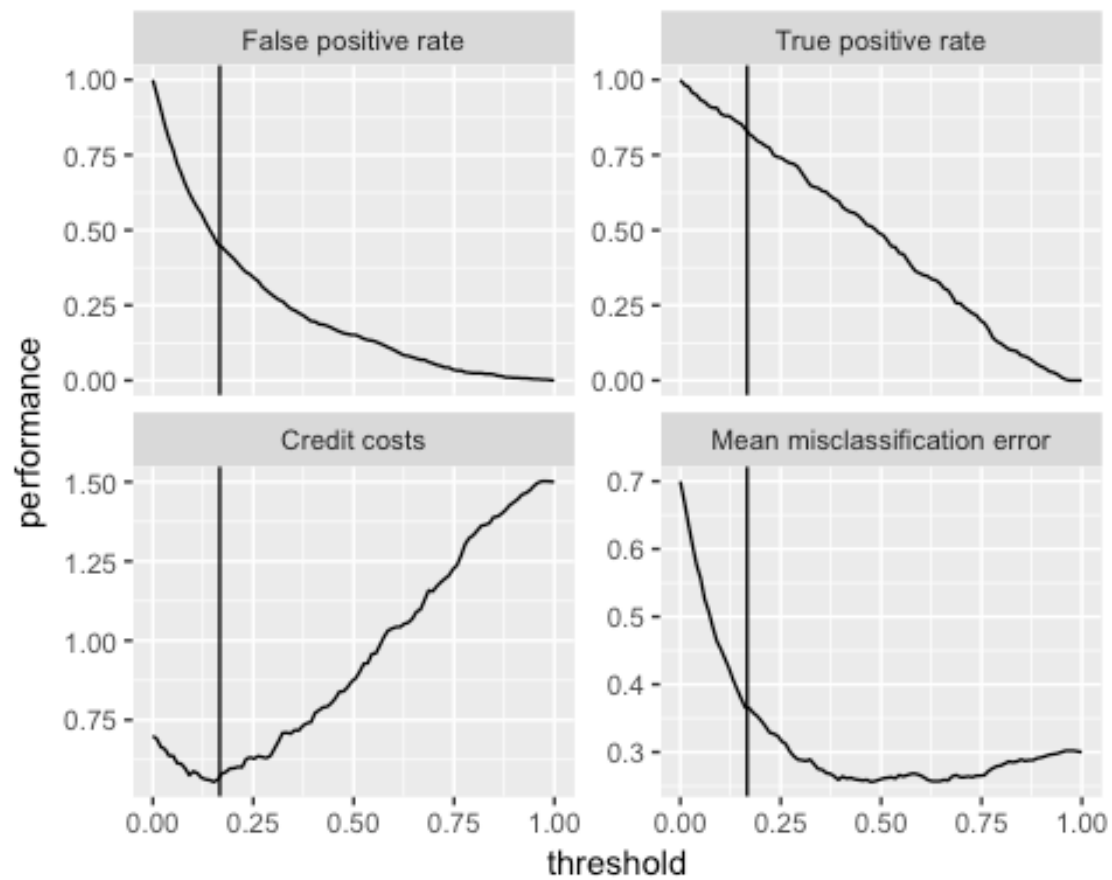
##           predicted
## true      Bad Good -err.-
##   Bad      249   51     51
##   Good      314  386    314
## -err.-      314   51    365

performance(Bpred.th, measures = list(credit.costs, mmce))


## credit.costs          mmce
##           0.569          0.365

Bir = generateThreshVsPerfData(Br, measures = list(fpr, tpr, credit.costs,
mmce))
plotThreshVsPerf(Bir, mark.th = th)

```



```
plotROCCurves(Bir)
```



```
performance(Br$pred,credit.costs)
```

```
## credit.costs
##          0.881
```

Neural Net Model

```
NNetltn = makeLearner("classif.nnet", predict.type = "prob",
  fix.factors.prediction = TRUE)
rin = makeResampleInstance("CV", iters = 5, task = credit.task, stratify=TRUE)
NNettr = resample(NNetltn, credit.task, rin, measures = list(credit.costs,
  mmce), show.info = FALSE)
```

```
## # weights: 265
## initial value 493.157921
## iter 10 value 370.961738
## iter 20 value 299.729609
## iter 30 value 272.215324
## iter 40 value 255.989264
## iter 50 value 253.281812
## iter 60 value 252.543626
## iter 70 value 251.555030
## iter 80 value 250.791717
```

```
## iter 90 value 250.715648
## iter 100 value 250.677780
## final value 250.677780
## stopped after 100 iterations
## # weights: 265
## initial value 518.390172
## iter 10 value 325.858908
## iter 20 value 238.530603
## iter 30 value 187.325546
## iter 40 value 167.850332
## iter 50 value 158.609335
## iter 60 value 152.482910
## iter 70 value 144.564783
## iter 80 value 142.580369
## iter 90 value 142.475667
## iter 100 value 142.464218
## final value 142.464218
## stopped after 100 iterations
## # weights: 265
## initial value 489.190637
## iter 10 value 395.997414
## iter 20 value 360.054362
## iter 30 value 334.448350
## iter 40 value 321.705984
## iter 50 value 319.144335
## iter 60 value 318.982250
## iter 70 value 318.978349
## iter 70 value 318.978349
## iter 70 value 318.978349
## final value 318.978349
## converged
## # weights: 265
## initial value 744.850551
## iter 10 value 442.605494
## iter 20 value 367.829756
## iter 30 value 354.029194
## iter 40 value 341.541711
## iter 50 value 327.922391
## iter 60 value 321.114031
## iter 70 value 320.596349
## iter 80 value 319.792041
## iter 90 value 319.371228
## iter 100 value 317.214438
## final value 317.214438
## stopped after 100 iterations
## # weights: 265
## initial value 502.617850
## iter 10 value 346.836484
## iter 20 value 304.668961
## iter 30 value 264.736920
```

```

## iter 40 value 238.870900
## iter 50 value 218.674601
## iter 60 value 217.311091
## iter 70 value 217.306192
## final value 217.306183
## converged

NNetr

## Resample Result
## Task: ml_credit_dataset
## Learner: classif.nnet
## Aggr perf: credit.costs.test.mean=0.8760000,mmce.test.mean=0.3160000
## Runtime: 0.95386

NNpred.th = setThreshold(NNetr$pred, threshold = th)
calculateConfusionMatrix(NNpred.th)


##           predicted
## true      Bad Good -err.-
## Bad      207  93    93
## Good     235 465    235
## -err.-   235  93    328

performance(NNpred.th, measures = list(credit.costs, mmce))


## credit.costs      mmce
##           0.700      0.328

NNr = generateThreshVsPerfData(NNetr, measures = list(fpr, tpr, credit.costs,
mmce))
plotThreshVsPerf(NNr, mark.th = th)

```



```
plotROCCurves(NNr)
```



```
performance(NNetr$pred,credit.costs)
```

```
## credit.costs
```

```
##          0.876
```