

German Credit Data Exploration_5

Dr. Prashant Mishra

3/27/2018

```
ml_credit_dataset <- read.csv("ml_credit_dataset.csv")
str(ml_credit_dataset)
```

```
## 'data.frame': 1000 obs. of 87 variables:
## $ CheckingAccountStatus.0.to.200 : int 0 1 0 0 0 0 0 1 0 1 ...
## $ CheckingAccountStatus.gt.200 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ CheckingAccountStatus.lt.0 : int 1 0 0 1 1 0 0 0 0 0 ...
## $ CheckingAccountStatus.none : int 0 0 1 0 0 1 1 0 1 0 ...
## $ Duration.0.to.6 : int 1 0 0 0 0 0 0 0 0 0 ...
## $ Duration.6.to.12 : int 0 0 1 0 0 0 0 0 1 0 ...
## $ Duration.12.to.18 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Duration.18.to.24 : int 0 0 0 0 1 0 1 0 0 0 ...
## $ Duration.24.to.30 : int 0 0 0 0 0 0 0 0 0 1 ...
## $ Duration.30.to.36 : int 0 0 0 0 0 1 0 1 0 0 ...
## $ Duration.36.to.42 : int 0 0 0 1 0 0 0 0 0 0 ...
## $ Duration.42.to.48 : int 0 1 0 0 0 0 0 0 0 0 ...
## $ Duration.48.to.54 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Duration.54.to.60 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Duration.66.to.72 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ CreditHistory.Critical : int 1 0 1 0 0 0 0 0 0 1 ...
## $ CreditHistory.Delay : int 0 0 0 0 1 0 0 0 0 0 ...
## $ CreditHistory.NoCredit.AllPaid : int 0 0 0 0 0 0 0 0 0 0 ...
## $ CreditHistory.PaidDuly : int 0 1 0 1 0 1 1 1 1 0 ...
## $ CreditHistory.ThisBank.AllPaid : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Purpose.Business : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Purpose.DomesticAppliance : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Purpose.Education : int 0 0 1 0 0 1 0 0 0 0 ...
## $ Purpose.Furniture.Equipment : int 0 0 0 1 0 0 1 0 0 0 ...
## $ Purpose.NewCar : int 0 0 0 0 1 0 0 0 0 1 ...
## $ Purpose.Others : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Purpose.Radio.Television : int 1 1 0 0 0 0 0 0 1 0 ...
## $ Purpose.Repairs : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Purpose.Retaining : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Purpose.UsedCar : int 0 0 0 0 0 0 0 1 0 0 ...
## $ SavingsAccountBonds.100.to.500 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ SavingsAccountBonds.500.to.1000 : int 0 0 0 0 0 0 1 0 0 0 ...
## $ SavingsAccountBonds.gt.1000 : int 0 0 0 0 0 0 0 0 1 0 ...
## $ SavingsAccountBonds.lt.100 : int 0 1 1 1 1 0 0 1 0 1 ...
## $ SavingsAccountBonds.Unknown : int 1 0 0 0 0 1 0 0 0 0 ...
## $ EmploymentDuration.0.to.1 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ EmploymentDuration.1.to.4 : int 0 1 0 0 1 1 0 1 0 0 ...
## $ EmploymentDuration.4.to.7 : int 0 0 1 1 0 0 0 0 1 0 ...
## $ EmploymentDuration.gt.7 : int 1 0 0 0 0 0 1 0 0 0 ...
## $ EmploymentDuration.Unemployed : int 0 0 0 0 0 0 0 0 0 1 ...
## $ InstallmentRatePercentage.1 : int 0 0 0 0 0 0 0 0 0 0 ...
## $ InstallmentRatePercentage.2 : int 0 1 1 1 0 1 0 1 1 0 ...
## $ InstallmentRatePercentage.3 : int 0 0 0 0 1 0 1 0 0 0 ...
```

```

## $ InstallmentRatePercentage.4      : int  1 0 0 0 0 0 0 0 0 1 ...
## $ Personal.Female.NotSingle        : int  0 1 0 0 0 0 0 0 0 0 ...
## $ Personal.Male.Divorced.Seperated : int  0 0 0 0 0 0 0 0 1 0 ...
## $ Personal.Male.Married.Widowed    : int  0 0 0 0 0 0 0 0 0 1 ...
## $ Personal.Male.Single              : int  1 0 1 1 1 1 1 1 0 0 ...
## $ OtherDebtorsGuarantors.CoApplicant : int  0 0 0 0 0 0 0 0 0 0 ...
## $ OtherDebtorsGuarantors.Guarantor  : int  0 0 0 1 0 0 0 0 0 0 ...
## $ OtherDebtorsGuarantors.None       : int  1 1 1 0 1 1 1 1 1 1 ...
## $ ResidenceDuration.1               : int  0 0 0 0 0 0 0 0 0 0 ...
## $ ResidenceDuration.2               : int  0 1 0 0 0 0 0 1 0 1 ...
## $ ResidenceDuration.3               : int  0 0 1 0 0 0 0 0 0 0 ...
## $ ResidenceDuration.4               : int  1 0 0 1 1 1 1 0 1 0 ...
## $ Property.CarOther                 : int  0 0 0 0 0 0 0 1 0 1 ...
## $ Property.Insurance                 : int  0 0 0 1 0 0 1 0 0 0 ...
## $ Property.RealEstate                : int  1 1 1 0 0 0 0 0 1 0 ...
## $ Property.Unknown                   : int  0 0 0 0 1 1 0 0 0 0 ...
## $ Age.18.to.24                      : int  0 1 0 0 0 0 0 0 0 0 ...
## $ Age.24.to.30                      : int  0 0 0 0 0 0 0 0 0 1 ...
## $ Age.30.to.36                      : int  0 0 0 0 0 1 0 1 0 0 ...
## $ Age.36.to.42                      : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Age.42.to.48                      : int  0 0 0 1 0 0 0 0 0 0 ...
## $ Age.48.to.54                      : int  0 0 1 0 1 0 1 0 0 0 ...
## $ Age.54.to.60                      : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Age.60.to.66                      : int  0 0 0 0 0 0 0 0 1 0 ...
## $ Age.66.to.72                      : int  1 0 0 0 0 0 0 0 0 0 ...
## $ Age.72.to.78                      : int  0 0 0 0 0 0 0 0 0 0 ...
## $ OtherInstallmentPlans.Bank         : int  0 0 0 0 0 0 0 0 0 0 ...
## $ OtherInstallmentPlans.None         : int  1 1 1 1 1 1 1 1 1 1 ...
## $ OtherInstallmentPlans.Stores       : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Housing.ForFree                   : int  0 0 0 1 1 1 0 0 0 0 ...
## $ Housing.Own                       : int  1 1 1 0 0 0 1 0 1 1 ...
## $ Housing.Rent                       : int  0 0 0 0 0 0 0 1 0 0 ...
## $ NumberExistingCredits.1            : int  0 1 1 1 0 1 1 1 1 0 ...
## $ NumberExistingCredits.2            : int  1 0 0 0 1 0 0 0 0 1 ...
## $ NumberExistingCredits.3            : int  0 0 0 0 0 0 0 0 0 0 ...
## $ NumberExistingCredits.4            : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Job.Management.SelfEmp.HighlyQualified: int  0 0 0 0 0 0 0 1 0 1 ...
## $ Job.SkilledEmployee                : int  1 1 0 1 1 0 1 0 0 0 ...
## $ Job.UnemployedUnskilled            : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Job.UnskilledResident              : int  0 0 1 0 0 1 0 0 1 0 ...
## $ NumberPeopleMaintenance            : int  1 1 2 2 2 2 1 1 1 1 ...
## $ Telephone                          : int  1 0 0 0 0 1 0 1 0 0 ...
## $ ForeignWorker                      : int  1 1 1 1 1 1 1 1 1 1 ...
## $ Class                              : Factor w/ 2 levels "Bad","Good": 2 1 2 2 1 2 2 2 2 1 ...

```

Making a Machine Learning task using mlr

```
library(mlr)
```

```
## Loading required package: ParamHelpers
```

```
credit.task = makeClassifTask(data = ml_credit_dataset, target = "Class")
credit.task = removeConstantFeatures(credit.task)
credit.task
```

```
## Supervised task: ml_credit_dataset
## Type: classif
## Target: Class
## Observations: 1000
## Features:
##      numerics      factors      ordered functionals
##           86           0           0           0
## Missings: FALSE
## Has weights: FALSE
## Has blocking: FALSE
## Has coordinates: FALSE
## Classes: 2
##      Bad Good
##      300 700
## Positive class: Bad
```

Cost Matrix for German Credit Data

```
costs = matrix(c(0, 1, 5, 0), 2)
colnames(costs) = rownames(costs) = getTaskClassLevels(credit.task)
costs
```

```
##      Bad Good
## Bad    0    5
## Good   1    0
```

Calculate the theoretical threshold for the positive class: Since $c(+1,+1)=c(-1,-1)=0$

```
th = costs[2,1]/(costs[2,1] + costs[1,2])
th
```

```
## [1] 0.1666667
```

Creating a cost measure

In order to calculate the average costs over the entire data set we first need to create a new performance Measure. This can be done through function `makeCostMeasure`. It is expected that the rows of the cost matrix indicate true and the columns predicted class labels.

```
credit.costs = makeCostMeasure(id = "credit.costs", name = "Credit costs", costs = costs,
  best = 0, worst = 5)
credit.costs
```

```
## Name: Credit costs
## Performance measure: credit.costs
## Properties: classif,classif.multi,req.pred,req.truth,predtype.response,predtype.prob
## Minimize: TRUE
## Best: 0; Worst: 5
## Aggregated by: test.mean
## Arguments: costs=<matrix>, combine=<function>
## Note:
```

2. Rebalancing

-In order to minimize the average costs, observations from the less costly class should be given higher importance during training.

-This can be achieved by weighting the classes, provided that the learner under consideration has a 'class weights' or an 'observation weights' argument.

i. Weighing

Just as theoretical thresholds, theoretical weights can be calculated from the cost matrix. If t indicates the target threshold and t_0 the original threshold for the positive class the proportion of observations in the positive class has to be multiplied by

$$w = \frac{1-t}{t} \frac{t_0}{1-t_0}$$

for our case: Weight for positive class corresponding to theoretical threshold

```
w = (1 - th)/th
w
```

```
## [1] 5
```

Assigning theoretical weight : for learner that support observation weights

-A unified and convenient way to assign class weights to a Learner (and tune them) is provided by function `makeWeightedClassesWrapper`.

-The class weights are specified using argument `wcw.weight`

-For learners that support observation weights a suitable weight vector is then generated internally during training or resampling.

```
wlrn = makeLearner("classif.multinom", trace = FALSE)
wlrn = makeWeightedClassesWrapper(wlrn, wcw.weight = w)
wlrn

## Learner weightedclasses.classif.multinom from package nnet
## Type: classif
## Name: ; Short name:
## Class: WeightedClassesWrapper
## Properties: twoclass,multiclass,numerics,factors,prob
## Predict-Type: response
## Hyperparameters: trace=FALSE,wcw.weight=5

rin = makeResampleInstance("CV", iters = 5, task = credit.task, stratify = TRUE)
wr = resample(wlrn, credit.task, rin, measures = list(credit.costs, mmce), show.info = FALSE)
wr

## Resample Result
## Task: ml_credit_dataset
## Learner: weightedclasses.classif.multinom
## Aggr perf: credit.costs.test.mean=0.5970000,mmce.test.mean=0.3730000
## Runtime: 0.674418
```

Assigning theoretical weight : for learner that support class weights

- If the learner can deal with class weights, the weights are basically passed on to the appropriate learner parameter.
- The advantage of using the wrapper in this case is the unified way to specify the class weights.
- For classification methods like “`classif.multinom`” that support class weights you can pass them directly.

```
lrn = makeWeightedClassesWrapper("classif.multinom", wcw.weight = w)
r = resample(lrn, credit.task, rin, measures = list(credit.costs, mmce), show.info = FALSE)
```

```
## # weights:  88 (87 variable)
## initial  value 1219.939038
## iter   10 value 796.373981
## iter   20 value 780.383943
## iter   30 value 778.920195
## iter   40 value 778.722353
## iter   50 value 778.606105
## iter   60 value 778.525562
## iter   70 value 778.513262
## iter   80 value 778.509883
## final   value 778.509795
## converged
## # weights:  88 (87 variable)
## initial  value 1219.939038
## iter   10 value 793.920599
## iter   20 value 778.699598
## iter   30 value 777.356296
## iter   40 value 777.263258
## iter   50 value 777.257323
## final   value 777.257050
## converged
## # weights:  88 (87 variable)
## initial  value 1219.939038
## iter   10 value 768.220550
## iter   20 value 748.106962
## iter   30 value 746.440508
## iter   40 value 746.241258
## iter   50 value 746.191259
## iter   60 value 746.139957
## iter   70 value 746.119464
## iter   80 value 746.115195
## final   value 746.115066
## converged
## # weights:  88 (87 variable)
## initial  value 1219.939038
## iter   10 value 791.473679
## iter   20 value 776.291122
## iter   30 value 774.558735
## iter   40 value 774.410797
## iter   50 value 774.338176
## iter   60 value 774.293733
## iter   70 value 774.282060
## iter   80 value 774.278767
```

```
## final value 774.278727
## converged
## # weights: 88 (87 variable)
## initial value 1219.939038
## iter 10 value 766.781471
## iter 20 value 744.967491
## iter 30 value 742.961794
## iter 40 value 742.541777
## iter 50 value 742.364527
## iter 60 value 742.250076
## iter 70 value 742.224207
## iter 80 value 742.216956
## final value 742.216728
## converged
```

```
r
```

```
## Resample Result
## Task: ml_credit_dataset
## Learner: weightedclasses.classif.multinom
## Aggr perf: credit.costs.test.mean=0.5970000,mmce.test.mean=0.3730000
## Runtime: 0.810432
```

Tuning the weight

-Just like the theoretical threshold, the theoretical weights may not always be suitable, therefore you can tune the weight for the positive class.

-Calculating the theoretical weight beforehand may help to narrow down the search interval.

```
lrn = makeLearner("classif.multinom", trace = FALSE)
lrn = makeWeightedClassesWrapper(lrn)
ps = makeParamSet(makeDiscreteParam("wcw.weight", seq(4, 12, 0.5)))
ctrl = makeTuneControlGrid()
tune.wcw.res = tuneParams(lrn, credit.task, resampling = rin, par.set = ps,
  measures = list(credit.costs, mmce), control = ctrl, show.info = FALSE)
tune.wcw.res
```

```
## Tune result:
## Op. pars: wcw.weight=5.5
## credit.costs.test.mean=0.5960000,mmce.test.mean=0.3840000
```

```
as.data.frame(tune.wcw.res$opt.path)[1:3]
```

	wcw.weight	credit.costs.test.mean	mmce.test.mean
## 1	4	0.608	0.348
## 2	4.5	0.604	0.364
## 3	5	0.597	0.373
## 4	5.5	0.596	0.384
## 5	6	0.605	0.401
## 6	6.5	0.615	0.411
## 7	7	0.618	0.418
## 8	7.5	0.625	0.429
## 9	8	0.630	0.438
## 10	8.5	0.628	0.440
## 11	9	0.625	0.445

## 12	9.5	0.604	0.444
## 13	10	0.600	0.448
## 14	10.5	0.607	0.455
## 15	11	0.613	0.461
## 16	11.5	0.621	0.469
## 17	12	0.626	0.474

ii. Over- and undersampling

-If the Learner supports neither observation nor class weights the proportions of the classes in the training data can be changed by over- or undersampling.

-In the GermanCredit data set the positive class Bad should receive a theoretical weight of $w = (1 - th)/th = 5$. This can be achieved by oversampling class Bad with a rate of 5 or by undersampling class Good with a rate of $1/5$ (using functions oversample or undersample).

logistic model

```
credit.task.over = oversample(credit.task, rate = w, cl = "Bad")
logisticlrn = makeLearner("classif.multinom", trace = FALSE)
logisticmod = mlr::train(logisticlrn, credit.task.over)
logisticpred = predict(logisticmod, task = credit.task)
performance(logisticpred, measures = list(credit.costs, mmce))
```

```
## credit.costs      mmce
##           0.443      0.323
```

Rpart model

```
credit.task.over = oversample(credit.task, rate = w, cl = "Bad")
rpartlrn = makeLearner("classif.rpart")
rpartmod = mlr::train(rpartlrn, credit.task.over)
rpartpred = predict(rpartmod, task = credit.task)
performance(rpartpred, measures = list(credit.costs, mmce))
```

```
## credit.costs      mmce
##           0.460      0.408
```

Resample data to get appropriate performance

-We usually prefer resampled performance values, but simply calling resample on the oversampled task does not work since predictions have to be based on the original task.

-The solution is to create a wrapped Learner via function makeOversampleWrapper.

-Internally, oversample is called before training, but predictions are done on the original data.

logistic model

```
logicallyrn = makeLearner("classif.multinom", trace = FALSE)
logicallyrn = makeOversampleWrapper(logicallyrn, osw.rate = w, osw.cl = "Bad")
logicallyrn
```

```
## Learner classif.multinom.oversampled from package mlr,nnet
## Type: classif
## Name: ; Short name:
```

```

## Class: OversampleWrapper
## Properties: numerics,factors,weights,prob,twoclass,multiclass
## Predict-Type: response
## Hyperparameters: trace=FALSE,osw.rate=5,osw.cl=Bad

lr = resample(logicallrn, credit.task, rin, measures = list(credit.costs, mmce), show.info = FALSE)
lr

## Resample Result
## Task: ml_credit_dataset
## Learner: classif.multinom.oversampled
## Aggr perf: credit.costs.test.mean=0.6080000,mmce.test.mean=0.3800000
## Runtime: 1.15702

Rpart model

rpartlrn = makeLearner("classif.rpart")
rpartlrn = makeOversampleWrapper(rpartlrn, osw.rate = w, osw.cl = "Bad")
rpartlrn

## Learner classif.rpart.oversampled from package mlr,rpart
## Type: classif
## Name: ; Short name:
## Class: OversampleWrapper
## Properties: numerics,factors,ordered,missings,weights,prob,twoclass,multiclass,featimp
## Predict-Type: response
## Hyperparameters: xval=0,osw.rate=5,osw.cl=Bad

rr = resample(logicallrn, credit.task, rin, measures = list(credit.costs, mmce), show.info = FALSE)
rr

## Resample Result
## Task: ml_credit_dataset
## Learner: classif.multinom.oversampled
## Aggr perf: credit.costs.test.mean=0.6110000,mmce.test.mean=0.3750000
## Runtime: 1.10054

```

Tuning the oversample rate

-Of course, we can also tune the oversampling rate. For this purpose we again have to create an OversampleWrapper. Optimal values for parameter osw.rate can be obtained using function tuneParams.

logistic model

```

logicallrn = makeLearner("classif.multinom", trace = FALSE)
logicallrn = makeOversampleWrapper(logicallrn, osw.cl = "Bad")
logicalps = makeParamSet(makeDiscreteParam("osw.rate", seq(3, 8, 0.25)))
logicalctrl = makeTuneControlGrid()
logicaltune.osw.res = tuneParams(logicallrn, credit.task, rin, par.set = logicalps, measures = list(credit.costs, mmce),
  control = logicalctrl, show.info = FALSE)
logicaltune.osw.res

## Tune result:
## Op. pars: osw.rate=5
## credit.costs.test.mean=0.5820000,mmce.test.mean=0.3660000

```

Rpart model


```

rpartlrn = makeLearner("classif.rpart")
rpartlrn = makeOversampleWrapper(rpartlrn, osw.cl = "Bad")
rpartps = makeParamSet(makeDiscreteParam("osw.rate", seq(3, 8, 0.25)))
rpartctrl = makeTuneControlGrid()
rparttune.osw.res = tuneParams(rpartlrn, credit.task, rin, par.set = rpartps, measures = list(credit.co
  control = rpartctrl, show.info = FALSE)
rparttune.osw.res

```

Tune result:

Op. pars: osw.rate=7.75

credit.costs.test.mean=0.5660000,mmce.test.mean=0.4500000